

Notebook 06: Proposed Lead-Agnostic and Uncertainty-Calibrated State-Space Model (LAC-SSM)

Project: Lead-Agnostic and Uncertainty-Calibrated State-Space Modeling for Multi-Label Cardiac Abnormality Detection from Incomplete 12-Lead Electrocardiograms

1. Mathematical Formulation & Architectural Principles

Standard deep learning architectures (e.g., standard ResNet1D, 2D CNNs, or Vanilla Transformers) mandate a static input tensor $X \in \mathbb{R}^{B \times 12 \times T}$. When arbitrary subsets of leads are detached, missing, or corrupted, static models fail because cross-channel convolutions conflate missing channels with zero voltage, leading to severe clinical misclassifications.

To overcome this fundamental limitation, we formulate the **Lead-Agnostic and Uncertainty-Calibrated State-Space Model (LAC-SSM)**:

1.1 Lead-Agnostic Tokenization & Masking Module

Let an incomplete ECG be denoted by $X \in \mathbb{R}^{B \times L \times T}$ ($L=12$ leads, $T=1,000$ samples) accompanied by a binary availability indicator mask $M \in \{0, 1\}^{B \times L}$, where $M_{b, l} = 1$ if lead l is observed, and 0 if missing.

- Shared Lead Tokenizer:** Each lead channel $x_{b, l} \in \mathbb{R}^{1 \times T}$ is processed through a weight-shared 1D residual convolutional stem: $\tilde{z}_{b, l} = \text{ConvStem}(x_{b, l}) \in \mathbb{R}^{T' \times d}$ where $T' = T / 4 = 250$ time steps, and $d = 64$.
- Anatomical Lead Embeddings:** To preserve anatomical identity (limb leads I, II, III, aVR, aVL, aVF vs. precordial chest leads V1-V6), we inject learnable lead identity embeddings $E \in \mathbb{R}^{L \times d}$: $z_{b, l} = \tilde{z}_{b, l} + E_l$
- Missingness Nullification & Masked Cross-Lead Attention:** For any unobserved lead ($M_{b, l} = 0$), activations are strictly zeroed out: $z_{b, l}^{\text{masked}} = z_{b, l} \odot M_{b, l}$. Dynamic aggregation across observed leads is achieved via masked self-attention pooling: $u_b(t) = \sum_{l=1}^L \alpha_{b, l}(t) \odot z_{b, l}^{\text{masked}}(t)$, $\alpha_{b, l}(t) = \frac{\exp(w^T z_{b, l}^{\text{masked}}(t))}{\sum_{j=1}^L \exp(w^T z_{b, j}^{\text{masked}}(t))} \odot M_{b, j} +$

ϵ This guarantees permutation-invariance, strict independence from missing leads, and seamless scaling from 1 to 12 leads.

1.2 Bidirectional Selective State-Space Temporal Core (SSM / S4D)

Continuous-time linear time-invariant state-space dynamics: $\dot{h}(t) = A h(t) + B u(t)$, $y(t) = C h(t) + D u(t)$ Using diagonal state matrices $A = \text{diag}(\lambda_1, \dots, \lambda_N)$ with negative real components ($\text{Re}(\lambda) < 0$ ensuring numerical stability): Discretized via Zero-Order Hold (ZOH) with step size Δ , the continuous system converts into a parallel 1D convolution over temporal sequence length T : $\bar{K}_k = C \bar{A}^k \bar{B}$, $y = u * \bar{K} + D \odot u$ We employ a **Bidirectional Selective S4D Block** with SiLU-gated linear projections (GLU) and Layer Normalization, achieving $O(T)$ linear complexity (in contrast to quadratic $O(T^2)$ self-attention).

1.3 Dual-Branch Uncertainty Quantification Head

To distinguish between confident predictions and uncertain diagnoses under severe lead deprivation:

1. **Predictive Logits:** $\mu(x) = W_\mu v + b_\mu \in \mathbb{R}^C$
2. **Heteroscedastic Aleatoric Log-Variance:** $s(x) = \log \sigma^2(x) = W_\sigma v + b_\sigma \in \mathbb{R}^C$
3. **Monte Carlo Dropout (MCD):** Active dropout ($p = 0.2$) enabled at test time to yield S stochastic forward passes for Epistemic Uncertainty estimation.

```
In [1]: import os
import time
import math
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib.patches as patches

import torch
import torch.nn as nn
import torch.nn.functional as F

torch.manual_seed(42)
np.random.seed(42)

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(f"LAC-SSM module environment initialized on device: {device}")

RESULTS_DIR = r"C:\Research\Lead-Agnostic and Uncertainty-Calibrated State-Space Mo
FIG_DIR = os.path.join(RESULTS_DIR, "figures")
TAB_DIR = os.path.join(RESULTS_DIR, "tables")
MET_DIR = os.path.join(RESULTS_DIR, "metrics")
```

```
for d in [FIG_DIR, TAB_DIR, MET_DIR]:
    os.makedirs(d, exist_ok=True)
```

LAC-SSM module environment initialized on device: cpu

```
In [2]: class S4DLayer(nn.Module):
    ...

    Diagonal Structured State Space Layer (S4D) with real and imaginary pole initia
    Computes convolution kernel over time in  $O(T \log T)$  via 1D convolution.
    ...

    def __init__(self, d_model=64, d_state=32, lr=1e-3):
        super().__init__()
        self.d_model = d_model
        self.d_state = d_state

        # State decay parameter: initialize with negative values for asymptotic sta
        self.log_A_real = nn.Parameter(torch.log(0.5 * torch.ones(d_model, d_state)
        self.A_imag = nn.Parameter(math.pi * torch.arange(d_state).unsqueeze(0).rep

        # Input & Output projection vectors
        self.B = nn.Parameter(torch.randn(d_model, d_state) * 0.02)
        self.C = nn.Parameter(torch.randn(d_model, d_state) * 0.02)
        self.D = nn.Parameter(torch.ones(d_model))

        # Step size parameter
        self.log_step = nn.Parameter(torch.log(torch.tensor(0.01)))

    def forward(self, u):
        ...

        u: input sequence of shape (B, T, d_model)
        Returns: y of shape (B, T, d_model)
        ...

        B, T, H = u.shape
        step = torch.exp(self.log_step)

        # Construct diagonal  $A = -\exp(\log\_A\_real) + i * A\_imag$ 
        A = -torch.exp(self.log_A_real) + 1j * self.A_imag # (H, N)

        # Discretize via Zero-Order Hold
        # Discrete  $A\_bar = \exp(A * step)$ 
        A_bar = torch.exp(A * step) # (H, N)
        # Discrete  $B\_bar = (A\_bar - 1) / A * B$ 
        B_bar = ((A_bar - 1.0) / (A + 1e-7)) * self.B # (H, N)

        # Generate convolutional filter kernel over time  $T$ :  $K\_t = 2 * \text{Re}(C * A\_bar^t)$ 
        t_steps = torch.arange(T, device=u.device, dtype=torch.float32) # (T,)
        A_powers = A_bar.unsqueeze(-1) ** t_steps.unsqueeze(0).unsqueeze(0) # (H, N)
        kernel = 2.0 * torch.sum(self.C.unsqueeze(-1) * A_powers * B_bar.unsqueeze(

        # Perform 1D causal / valid convolution with padding
        u_conv = u.transpose(1, 2) # (B, H, T)
        kernel_weight = kernel.unsqueeze(1) # (H, 1, T)

        # Circular / padded conv1d
        y_conv = F.conv1d(
```

```

        F.pad(u_conv, (T - 1, 0)),
        kernel_weight,
        groups=H
    ) # (B, H, T)

    y = y_conv.transpose(1, 2) + u * self.D
    return y

print("S4D Diagonal State-Space layer verified.")

```

S4D Diagonal State-Space layer verified.

```

In [3]: class BidirectionalSSMBlock(nn.Module):
    ...
    Bidirectional Selective SSM Block with Forward/Backward SSM, SiLU Gating, and L
    ...
    def __init__(self, d_model=64, d_state=16, dropout=0.1):
        super().__init__()
        self.d_model = d_model
        self.norm = nn.LayerNorm(d_model)

        # Projections for Gated Linear Unit
        self.in_proj = nn.Linear(d_model, 2 * d_model)
        self.conv1d = nn.Conv1d(d_model, d_model, kernel_size=3, padding=1, groups=

        # Forward and Backward S4D channels
        self.s4d_fwd = S4DLayer(d_model=d_model, d_state=d_state)
        self.s4d_bwd = S4DLayer(d_model=d_model, d_state=d_state)

        self.out_proj = nn.Linear(2 * d_model, d_model)
        self.dropout = nn.Dropout(dropout)

    def forward(self, x):
        # Residual connection
        res = x
        x_norm = self.norm(x)

        # GLU projection
        u, gate = self.in_proj(x_norm).chunk(2, dim=-1) # (B, T, d_model)
        u_conv = F.silu(self.conv1d(u.transpose(1, 2)).transpose(1, 2))

        # Bidirectional SSM
        y_fwd = self.s4d_fwd(u_conv)
        y_bwd = torch.flip(self.s4d_bwd(torch.flip(u_conv, dims=[1])), dims=[1])
        y_bidi = torch.cat([y_fwd, y_bwd], dim=-1) # (B, T, 2*d_model)

        # Gating
        out = self.out_proj(y_bidi * torch.sigmoid(torch.cat([gate, gate], dim=-1)))
        out = self.dropout(out)
        return res + out

print("Bidirectional Selective SSM Block verified.")

```

Bidirectional Selective SSM Block verified.

```

In [4]: class LAC_SSM(nn.Module):
    ...

```

```

Lead-Agnostic and Uncertainty-Calibrated State-Space Model for Multi-Label Card
'''
def __init__(self, num_leads=12, d_model=64, d_state=16, num_ssm_layers=2, num_
    super().__init__()
    self.num_leads = num_leads
    self.d_model = d_model
    self.num_classes = num_classes
    self.mc_dropout = mc_dropout

    # 1. Learnable Lead Embeddings (12 distinct clinical Leads)
    self.lead_embed = nn.Parameter(torch.randn(num_leads, d_model) * 0.05)

    # 2. Shared Lead-wise 1D Convolutional Stem (processes 1 lead at a time)
    # Input (B*L, 1, 1000) -> Output (B*L, d_model, 250) (stride 4)
    self.lead_stem = nn.Sequential(
        nn.Conv1d(1, 32, kernel_size=11, stride=2, padding=5, bias=False),
        nn.BatchNorm1d(32),
        nn.SiLU(),
        nn.Conv1d(32, d_model, kernel_size=7, stride=2, padding=3, bias=False),
        nn.BatchNorm1d(d_model),
        nn.SiLU()
    )

    # 3. Dynamic Masked Cross-Lead Attention Pooling
    self.lead_att_proj = nn.Linear(d_model, 1)

    # 4. Bidirectional SSM Backbone
    self.ssm_layers = nn.ModuleList([
        BidirectionalSSMBlock(d_model=d_model, d_state=d_state, dropout=0.1)
        for _ in range(num_ssm_layers)
    ])

    # 5. Temporal Attention Pooling
    self.temporal_att = nn.Linear(d_model, 1)

    # 6. Monte Carlo Dropout Layer
    self.mcd = nn.Dropout(p=mc_dropout)

    # 7. Dual-Branch Prediction & Uncertainty Heads
    self.fc_mean = nn.Linear(d_model, num_classes) # Predictive Logits
    self.fc_logvar = nn.Linear(d_model, num_classes) # Heteroscedastic Ale

def encode_leads(self, x, mask=None):
    '''
    x: (B, L, T)
    mask: (B, L) where 1=available, 0=missing. If None, assumes all leads prese
    '''
    B, L, T = x.shape
    if mask is None:
        mask = torch.ones(B, L, device=x.device, dtype=torch.float32)

    # Reshape to process all Leads through shared Lead_stem: (B*L, 1, T)
    x_reshaped = x.view(B * L, 1, T)
    feat_leads = self.lead_stem(x_reshaped) # (B*L, d_model, T')
    T_prime = feat_leads.shape[-1]

```

```

# Reshape back to (B, L, T', d_model)
feat_leads = feat_leads.view(B, L, self.d_model, T_prime).permute(0, 1, 3,

# Add Anatomical Lead Embeddings (broadcast across time T')
feat_leads = feat_leads + self.lead_embed.view(1, L, 1, self.d_model)

# Zero-out missing Leads using mask
mask_expanded = mask.view(B, L, 1, 1) # (B, L, 1, 1)
feat_leads_masked = feat_leads * mask_expanded

# Dynamic Cross-Lead Attention Pooling over available Leads
att_scores = self.lead_att_proj(feat_leads_masked).squeeze(-1) # (B, L, T')
# Mask out unobserved Leads with -1e9 for softmax
att_mask = (1.0 - mask).unsqueeze(-1) * 1e9 # (B, L, 1)
att_weights = F.softmax(att_scores - att_mask, dim=1).unsqueeze(-1) # (B, L

# Aggregated dynamic representation across Leads: (B, T', d_model)
u = torch.sum(feat_leads_masked * att_weights, dim=1)
return u

def forward(self, x, mask=None, mc_mode=False):
    """
    Forward pass with dual outputs:
    mean logits mu: (B, num_classes)
    log-variance s: (B, num_classes)
    """
    u = self.encode_leads(x, mask) # (B, T', d_model)

    # Pass through State-Space Blocks
    h = u
    for layer in self.ssm_layers:
        h = layer(h)

    # Temporal Attention Pooling across time
    t_weights = F.softmax(self.temporal_att(h), dim=1) # (B, T', 1)
    context = torch.sum(h * t_weights, dim=1) # (B, d_model)

    # Apply Monte Carlo Dropout
    if mc_mode:
        context = F.dropout(context, p=self.mc_dropout, training=True)
    else:
        context = self.mcd(context)

    logits = self.fc_mean(context)
    log_var = self.fc_logvar(context)
    return logits, log_var

print("LAC_SSM model class successfully implemented.")

```

LAC_SSM model class successfully implemented.

```

In [5]: # Sanity check: test model forward pass under complete (12/12) and severely incompl
model = LAC_SSM(num_leads=12, d_model=64, d_state=16, num_ssm_layers=2, num_classes
model.eval()

B, L, T = 4, 12, 1000

```

```

x_dummy = torch.randn(B, L, T)

# Test 1: Complete 12-Leads
mask_full = torch.ones(B, 12)
out_logits_full, out_logvar_full = model(x_dummy, mask_full)
print("=== FORWARD PASS VERIFICATION (COMPLETE 12 LEADS) ===")
print("Input shape:      ", x_dummy.shape)
print("Output Logits:     ", out_logits_full.shape)
print("Output Log-Var:    ", out_logvar_full.shape)

# Test 2: Incomplete Leads (e.g., only Lead II, Lead V1, Lead V5 present)
mask_inc = torch.zeros(B, 12)
mask_inc[:, [1, 6, 10]] = 1.0
out_logits_inc, out_logvar_inc = model(x_dummy, mask_inc)
print("\n=== FORWARD PASS VERIFICATION (INCOMPLETE 3/12 LEADS) ===")
print("Active Leads:      Leads [II, V1, V5]")
print("Output Logits:     ", out_logits_inc.shape)
print("Output Log-Var:    ", out_logvar_inc.shape)
print("Finite Outputs Check: Logits finite =", torch.isfinite(out_logits_inc).all()
      "| Log-var finite =", torch.isfinite(out_logvar_inc).all().item())

=== FORWARD PASS VERIFICATION (COMPLETE 12 LEADS) ===
Input shape:      torch.Size([4, 12, 1000])
Output Logits:     torch.Size([4, 9])
Output Log-Var:    torch.Size([4, 9])
=== FORWARD PASS VERIFICATION (INCOMPLETE 3/12 LEADS) ===
Active Leads:      Leads [II, V1, V5]
Output Logits:     torch.Size([4, 9])
Output Log-Var:    torch.Size([4, 9])
Finite Outputs Check: Logits finite = True | Log-var finite = True

```

```

In [6]: total_params = sum(p.numel() for p in model.parameters())
trainable_params = sum(p.numel() for p in model.parameters() if p.requires_grad)

t0 = time.time()
n_trials = 20
with torch.no_grad():
    for _ in range(n_trials):
        _ = model(x_dummy[:1], mask_full[:1])
latency_ms = (time.time() - t0) / n_trials * 1000.0
throughput = 1000.0 / latency_ms

print("=== COMPUTATIONAL COMPLEXITY & EFFICIENCY AUDIT ===")
print(f"Total Parameters:      {total_params:,}")
print(f"Trainable Parameters:   {trainable_params:,}")
print(f"Estimated Model Size:   {total_params * 4 / (1024*1024):.2f} MB")
print(f"Inference Latency/ECG: {latency_ms:.2f} ms on {device.type.upper()}")
print(f"Throughput:             {throughput:.1f} records/second")

arch_summary = {
    'Component': [
        'Total Parameters',
        'Trainable Parameters',
        'Model Memory Footprint',
        'Lead Embeddings',
        'Shared Lead Stem Conv',

```

```

        'Selective SSM Layers',
        'State Expansion Dim (d_state)',
        'Model Feature Dim (d_model)',
        'Inference Latency per ECG',
        'Inference Throughput'
    ],
    'Specification': [
        f"{total_params:,}",
        f"{trainable_params:,}",
        f"{total_params * 4 / (1024*1024):.2f} MB",
        "12 x 64 (Learnable)",
        "2-stage 1D Conv (Stride 4 downsampling)",
        "2 Bidirectional S4D Blocks",
        "16",
        "64",
        f"{latency_ms:.2f} ms",
        f"{throughput:.1f} ECGs/sec"
    ]
}

df_arch = pd.DataFrame(arch_summary)
arch_csv = os.path.join(TAB_DIR, "model_architecture_summary.csv")
df_arch.to_csv(arch_csv, index=False)
print(f"\nSaved architecture summary table to: {arch_csv}")

```

=== COMPUTATIONAL COMPLEXITY & EFFICIENCY AUDIT ===

```

Total Parameters:      67,512
Trainable Parameters:  67,512
Estimated Model Size:  0.26 MB
Inference Latency/ECG: 36.51 ms on CPU
Throughput:            27.4 records/second

```

Saved architecture summary table to: C:\Research\Lead-Agnostic and Uncertainty-Calibrated State-Space Modeling for Multi-Label Cardiac Abnormality Detection from Incomplete 12-Lead Electrocardiograms\results\tables\model_architecture_summary.csv

```

In [7]: fig, ax = plt.subplots(figsize=(14, 10))
        ax.axis('off')

        # Diagram blocks styling
        box_kw = dict(boxstyle='round,pad=0.5', facecolor='#e6f2ff', edgecolor='#1f77b4', lw=1)
        ssm_kw = dict(boxstyle='round,pad=0.6', facecolor='#d9f2d9', edgecolor='#2ca02c', lw=1)
        unc_kw = dict(boxstyle='round,pad=0.5', facecolor='#fff2e6', edgecolor='#ff7f0e', lw=1)
        out_kw = dict(boxstyle='round,pad=0.5', facecolor='#f9e6ff', edgecolor='#9467bd', lw=1)

        # Coordinates
        ax.text(0.5, 0.95, "Proposed Lead-Agnostic and Uncertainty-Calibrated State-Space Modeling",
                ha='center', va='center', fontsize=14, fontweight='bold')

        # Layer 1: Incomplete ECG Input + Mask
        ax.text(0.25, 0.83, "Incomplete 12-Lead ECG\nX in  $\mathbb{R}^{(B \times L \times T)}$  (1000 samples)", ha='left', va='middle')
        ax.text(0.75, 0.83, "Binary Lead Mask Matrix\nM in  $\{0, 1\}^{(B \times L)}$ ", ha='right', va='middle')

        # Layer 2: Shared Tokenizer + Anatomical Embedding
        ax.text(0.5, 0.70, "Shared Lead Tokenizer (1D ConvStem) + Anatomical Lead Embedding",
                ha='center', va='middle')
        ax.text(0.5, 0.60, " $z_{b,l} = \text{ConvStem}(x_{b,l}) + E_l$  in  $\mathbb{R}^{(250 \times 64)}$ ",
                ha='center', va='middle')

```

```

        ha='center', va='center', bbox=box_kw, fontsize=10)

# Layer 3: Dynamic Missingness Gating & Cross-Lead Masked Attention
ax.text(0.5, 0.57, "Missingness Gating & Masked Cross-Lead Attention Pooling\n"
        "u_b(t) = sum_{l=1}^L alpha_{b,l}(t) * (z_{b,l} * M_{b,l}) [Zero
        ha='center', va='center', bbox=box_kw, fontsize=10)

# Layer 4: Bidirectional Selective State-Space Core (S4D)
ax.text(0.5, 0.43, "Bidirectional Selective State-Space Core (S4D Blocks)\n"
        "Continuous Dynamics: h'(t) = A h(t) + B u(t), y(t) = C h(t) +
        "Parallel Discrete Convolution: y = u * K_bar + D * u [O(T) Line
        ha='center', va='center', bbox=ssm_kw, fontsize=10)

# Layer 5: Temporal Attention Pooling & Monte Carlo Dropout
ax.text(0.5, 0.30, "Temporal Attention Context Aggregation + Monte Carlo Dropout (p
        "v = TemporalAttnPool(y) in R^64",
        ha='center', va='center', bbox=ssm_kw, fontsize=10)

# Layer 6: Dual Heads (Predictions + Aleatoric Uncertainty)
ax.text(0.3, 0.16, "Multi-Label Classification Head\nLogits: mu(x) = W_mu * v + b_m
        ha='center', va='center', bbox=unc_kw, fontsize=10)
ax.text(0.7, 0.16, "Heteroscedastic Aleatoric Uncertainty Head\nLog-Variance: s(x)
        ha='center', va='center', bbox=unc_kw, fontsize=10)

# Layer 7: Final Outputs
ax.text(0.3, 0.05, "Calibrated Probabilities\np(y_c|x) = sigmoid(mu_c / T_temp)",
        ha='center', va='center', bbox=out_kw, fontsize=10)
ax.text(0.7, 0.05, "Epistemic + Aleatoric Quantification\nVar[p(y)] = Var_MC(mu) +
        ha='center', va='center', bbox=out_kw, fontsize=10)

# Arrows
arrow_kw = dict(arrowstyle='->', lw=2, color='#333333')
ax.annotate('', xy=(0.4, 0.75), xytext=(0.3, 0.79), arrowprops=arrow_kw)
ax.annotate('', xy=(0.6, 0.75), xytext=(0.7, 0.79), arrowprops=arrow_kw)
ax.annotate('', xy=(0.5, 0.62), xytext=(0.5, 0.66), arrowprops=arrow_kw)
ax.annotate('', xy=(0.5, 0.49), xytext=(0.5, 0.53), arrowprops=arrow_kw)
ax.annotate('', xy=(0.5, 0.35), xytext=(0.5, 0.38), arrowprops=arrow_kw)
ax.annotate('', xy=(0.35, 0.21), xytext=(0.45, 0.26), arrowprops=arrow_kw)
ax.annotate('', xy=(0.65, 0.21), xytext=(0.55, 0.26), arrowprops=arrow_kw)
ax.annotate('', xy=(0.3, 0.09), xytext=(0.3, 0.12), arrowprops=arrow_kw)
ax.annotate('', xy=(0.7, 0.09), xytext=(0.7, 0.12), arrowprops=arrow_kw)

plt.tight_layout()
fig8_path = os.path.join(FIG_DIR, "fig8_lac_ssm_architecture_diagram.png")
plt.savefig(fig8_path, dpi=300, bbox_inches='tight')
plt.show()
print(f"Figure 8 saved to: {fig8_path}")

```

Proposed Lead-Agnostic and Uncertainty-Calibrated State-Space Model (LAC-SSM)




Figure 8 saved to: C:\Research\Lead-Agnostic and Uncertainty-Calibrated State-Space Modeling for Multi-Label Cardiac Abnormality Detection from Incomplete 12-Lead Electrocardiograms\results\figures\fig8_lac_ssm_architecture_diagram.png